

深度学习基础作业 1-1: FNN 回归任务实验报告

学号: PB23000270 姓名: 李佩优

2026 年 3 月 29 日

摘要

本实验基于糖尿病疾病进展数据集, 通过构建前馈神经网络 (FNN) 回归模型, 系统探究了网络深度、学习率与激活函数对模型性能的影响。实验采用均方误差作为损失函数, Adam 优化器, 并分别对比了三种深度 (浅层 [32]、中层 [64,32]、深层 [128,64,32])、四种学习率 (0.1、0.01、0.001、0.0001) 以及五种激活函数 (ReLU、LeakyReLU、Swish、Sigmoid、Tanh) 的设置。实验结果表明: 网络深度适中的中层网络泛化能力稳定, 深层网络虽取得最低测试误差 (2829.05) 但存在轻微过拟合; 学习率 0.01 可获得最优测试性能 (2634.36), 过大或过小均影响收敛与泛化; 激活函数方面, ReLU、LeakyReLU 与 Swish 显著优于 Sigmoid 和 Tanh, 其中 ReLU 以最低测试损失 (2910.50) 表现最佳。上述发现验证了深度学习理论中关于超参数选择的基本原则, 并为实际回归任务中的模型调优提供了参考。

1 引言

深度学习在回归任务中应用广泛, 其性能受到网络结构、优化设置和激活函数等多种因素影响。本报告基于 Diabetes 数据集, 通过设计对比实验, 旨在深入理解这些因素对前馈神经网络 (FNN) 训练过程和泛化能力的作用机制, 并为模型调优提供实践指导。

2 实验设置

2.1 数据集

实验使用 sklearn 中的 Diabetes 数据集, 包含 442 个样本, 每个样本有 10 个数值型特征。数据集已预先标准化。实验按照 4:1 的比例将数据划分为训练集 (80%) 和测试集 (20%), 并在训练集中进一步划分出验证集 (占训练集的 20%)。批量大小统一设置为 32。

2.2 模型与训练配置

我们构建了通用的前馈神经网络模型，其结构为：输入层（10 维）→ 若干隐藏层（每层使用指定激活函数）→ 输出层（1 维）。除对比实验外，默认配置如下：

- 损失函数: 均方误差 (MSE)
- 优化器: Adam
- 学习率: 0.001
- 隐藏层配置: [64, 32] (2 层)
- 激活函数: ReLU
- 训练轮数: 200

2.3 对比实验设计

为探究各因素的影响，我们设计了以下三组对比实验：

1. **网络深度影响:** 固定学习率 0.001 和 ReLU 激活函数，分别训练浅层网络 [32]、中层网络 [64,32] 和深层网络 [128,64,32]。
2. **学习率影响:** 固定中层网络 [64,32] 和 ReLU 激活函数，分别设置学习率为 0.1, 0.01, 0.001, 0.0001。
3. **激活函数影响:** 固定中层网络 [64,32] 和学习率 0.001，分别使用 Sigmoid、Tanh、ReLU、LeakyReLU 和 Swish 作为激活函数。

3 实验结果与分析

3.1 网络深度的影响




图 1: 不同网络深度的训练与验证损失曲线对比

图1展示了三种深度网络的训练和验证损失随轮次的变化。从图中可以观察到：

- **收敛速度与性能:** 浅层网络（[32]）训练和验证损失下降缓慢，最终损失远高于其他两种结构，存在明显欠拟合；深层网络（[128,64,32]）收敛最快，训练损失最终降至最低；中层网络（[64,32]）收敛速度介于两者之间。
- **过拟合现象:** 深层网络的验证损失在约 60 轮后不再下降，反而略有上升，而训练损失持续降低，表明出现了轻微过拟合；中层网络的训练与验证损失基本同步下降，未观察到明显过拟合。
- **测试集表现:** 浅层网络测试损失为 7815.46，中层为 2931.17，深层为 2829.05。深层网络在测试集上取得了最低误差，但验证损失略高于训练损失，存在轻微过拟合；中层网络泛化能力良好，测试损失与验证损失接近；浅层网络欠拟合严重，测试损失最高。

3.2 学习率的影响




图 2: 不同学习率的训练与验证损失曲线对比

图2展示了不同学习率下的损失曲线。测试损失分别为：

- 学习率 0.1 (Deci) : 2803.57
- 学习率 0.01 (Centi) : 2634.36
- 学习率 0.001 (Milli) : 2940.30
- 学习率 0.0001 (Mico) : 22242.82

不同学习率的影响十分显著：

- **学习率过大 (0.1):** 损失曲线前期下降最快，但后期剧烈震荡，无法稳定收敛，表明步长过大导致参数在最优解附近“弹跳”。尽管如此，其测试损失仍优于 0.001，说明在震荡中仍有部分 epoch 获得了较好泛化能力。
- **学习率适中 (0.01):** 损失曲线下降较快且整体平稳，后期虽有轻微波动但未发散，最终测试损失为所有设置中最低 (2634.36)，表明该学习率在本次任务中综合表现最佳。
- **学习率较小 (0.001):** 损失曲线平滑下降，收敛稳定，但最终测试损失 (2940.30) 略高于 0.1 和 0.01，可能原因是验证集最佳模型出现在较早阶段，未充分探索参数空间。
- **学习率过小 (0.0001):** 损失下降极其缓慢，在 200 轮内远未收敛，测试损失极高，说明该学习率无法有效更新参数。

3.3 激活函数的影响




图 3: 不同激活函数的训练与验证损失曲线对比

图3对比了五种激活函数的表现，最终测试损失为：

- ReLU: 2910.50
- LeakyReLU: 2935.83
- Swish: 2917.51
- Sigmoid: 13317.26
- Tanh: 12662.09

分析如下：

- **Sigmoid 和 Tanh:** 收敛速度极慢，200 轮后训练和验证损失仍处于较高水平，且测试损失远高于其他函数。这是由于饱和区的梯度消失问题，导致深层网络参数更新困难。
- **ReLU, LeakyReLU 和 Swish:** 三者收敛速度均较快，训练和验证损失在 40 轮后快速下降并趋于稳定。其中 ReLU 取得最低测试损失（2910.50），Swish 略高于 ReLU（2917.51），LeakyReLU 稍差（2935.83）。ReLU 的非饱和特性有效缓解了梯度消失，其简洁性在本任务中表现最佳；Swish 作为平滑变体性能接近 ReLU；LeakyReLU 在负区间的微小梯度并未带来明显优势。

4 结论与总结

通过本次实验，我们对影响 FNN 回归任务的关键因素有了更深入的认识。主要结论如下：

1. **网络深度**需要平衡模型容量与泛化能力。浅层网络欠拟合严重，深层网络虽取得最低测试误差但存在轻微过拟合，中层网络泛化性能稳定。对于本数据集，深层网络（[128,64,32]）在测试集上表现最佳，但需注意过拟合风险；中层网络（[64,32]）是更稳健的选择。
2. **学习率**是影响训练稳定性和收敛速度的关键超参数。过大会导致震荡，过小则收敛缓慢。本次实验中，0.01 的学习率取得了最低测试误差，0.1 虽震荡但性能尚可，0.001 收敛稳定但测试误差略高，0.0001 未收敛。实际调优时应结合损失曲线选择能使模型充分收敛且泛化能力最佳的学习率。
3. **激活函数**的选择至关重要。ReLU 及其变体（LeakyReLU, Swish）因其缓解梯度消失的特性，在训练效率和最终性能上均显著优于 Sigmoid 和 Tanh。

本次实验不仅验证了深度学习理论中的经典结论，也为实际模型调优提供了有价值的参考。在后续工作中，可以考虑引入正则化技术（如 Dropout、权重衰减）以进一步缓解深层网络的过拟合问题。

附录：代码说明

主要代码文件结构如下：

- `main.py`: 实验入口，选择运行哪个实验。
- `data/load_data.py`: 数据加载与预处理。
- `models/fnn.py`: 前馈神经网络模型定义。
- `train/trainer.py`: 训练、验证、评估流程封装。
- `experiments/`: 存放三个对比实验的脚本（`depth_experiment.py`, `lr_experiment.py`, `activation_experiment.py`）。

运行实验后，会自动生成相应的损失曲线图片（如 `depth_comparison.png`）并保存在当前目录下，报告中的图片即来源于此。

深度实验终端输出：

```
===== Training shallow network =====
Epoch 20/200, Train Loss: 28139.7349, Val Loss: 33532.9666
Epoch 40/200, Train Loss: 27021.1772, Val Loss: 32255.7603
Epoch 60/200, Train Loss: 25275.4782, Val Loss: 30272.0382
Epoch 80/200, Train Loss: 23059.2626, Val Loss: 27743.1599
Epoch 100/200, Train Loss: 20541.3658, Val Loss: 24856.6461
```

Epoch 120/200, Train Loss: 17889.8364, Val Loss: 21788.9724
Epoch 140/200, Train Loss: 15273.3030, Val Loss: 18708.2339
Epoch 160/200, Train Loss: 12814.2690, Val Loss: 15768.7348
Epoch 180/200, Train Loss: 10619.4559, Val Loss: 13084.4434
Epoch 200/200, Train Loss: 8762.1188, Val Loss: 10737.7634

===== Training medium network =====

Epoch 20/200, Train Loss: 24765.8325, Val Loss: 29352.6453
Epoch 40/200, Train Loss: 7481.4825, Val Loss: 8605.5296
Epoch 60/200, Train Loss: 4123.7281, Val Loss: 3540.6026
Epoch 80/200, Train Loss: 3801.9459, Val Loss: 3235.7236
Epoch 100/200, Train Loss: 3570.5788, Val Loss: 3086.4268
Epoch 120/200, Train Loss: 3397.7061, Val Loss: 2980.6469
Epoch 140/200, Train Loss: 3263.7724, Val Loss: 2950.5164
Epoch 160/200, Train Loss: 3163.9107, Val Loss: 2911.4266
Epoch 180/200, Train Loss: 3093.4946, Val Loss: 2887.6490
Epoch 200/200, Train Loss: 3039.4399, Val Loss: 2892.4500

===== Training deep network =====

Epoch 20/200, Train Loss: 4445.6936, Val Loss: 3581.2905
Epoch 40/200, Train Loss: 3529.9736, Val Loss: 3048.1187
Epoch 60/200, Train Loss: 3136.6955, Val Loss: 2891.8095
Epoch 80/200, Train Loss: 2977.5898, Val Loss: 2895.4027
Epoch 100/200, Train Loss: 2918.1942, Val Loss: 2941.2795
Epoch 120/200, Train Loss: 2897.3428, Val Loss: 2973.4696
Epoch 140/200, Train Loss: 2881.8500, Val Loss: 2945.4812
Epoch 160/200, Train Loss: 2883.0421, Val Loss: 2915.8524
Epoch 180/200, Train Loss: 2865.3630, Val Loss: 2921.9911
Epoch 200/200, Train Loss: 2854.0432, Val Loss: 2914.1573
shallow test loss: 7815.4584
medium test loss: 2931.1737
deep test loss: 2829.0488

步长实验终端输出:

===== Training Deci network =====

Epoch 20/200, Train Loss: 3269.8549, Val Loss: 2979.2036
Epoch 40/200, Train Loss: 2754.2628, Val Loss: 2845.8914

Epoch 60/200, Train Loss: 2792.4984, Val Loss: 3690.0559
Epoch 80/200, Train Loss: 2852.1469, Val Loss: 3302.4515
Epoch 100/200, Train Loss: 2841.6267, Val Loss: 2873.5820
Epoch 120/200, Train Loss: 2648.9875, Val Loss: 3244.9714
Epoch 140/200, Train Loss: 2430.3948, Val Loss: 3206.5097
Epoch 160/200, Train Loss: 2421.9041, Val Loss: 3267.6666
Epoch 180/200, Train Loss: 2683.1864, Val Loss: 3892.2422
Epoch 200/200, Train Loss: 2451.0956, Val Loss: 2864.6637

===== Training Centi network =====

Epoch 20/200, Train Loss: 3115.1314, Val Loss: 2844.3065
Epoch 40/200, Train Loss: 3013.1535, Val Loss: 2774.4965
Epoch 60/200, Train Loss: 2911.2458, Val Loss: 2846.3074
Epoch 80/200, Train Loss: 2887.8727, Val Loss: 2888.6388
Epoch 100/200, Train Loss: 2849.0019, Val Loss: 2896.7128
Epoch 120/200, Train Loss: 2866.5126, Val Loss: 3012.4425
Epoch 140/200, Train Loss: 2825.7127, Val Loss: 2816.7726
Epoch 160/200, Train Loss: 2805.7143, Val Loss: 2973.5132
Epoch 180/200, Train Loss: 2779.3247, Val Loss: 2859.6492
Epoch 200/200, Train Loss: 2766.8575, Val Loss: 2990.0087

===== Training Milli network =====

Epoch 20/200, Train Loss: 24809.3014, Val Loss: 29451.9156
Epoch 40/200, Train Loss: 8087.2564, Val Loss: 9373.9606
Epoch 60/200, Train Loss: 4170.7511, Val Loss: 3616.5900
Epoch 80/200, Train Loss: 3839.8955, Val Loss: 3261.0513
Epoch 100/200, Train Loss: 3606.3579, Val Loss: 3107.3309
Epoch 120/200, Train Loss: 3428.0162, Val Loss: 3005.0018
Epoch 140/200, Train Loss: 3291.6069, Val Loss: 2957.9178
Epoch 160/200, Train Loss: 3186.4542, Val Loss: 2910.2736
Epoch 180/200, Train Loss: 3109.5299, Val Loss: 2899.6381
Epoch 200/200, Train Loss: 3053.1374, Val Loss: 2886.5619

===== Training Mico network =====

Epoch 20/200, Train Loss: 28511.5848, Val Loss: 33971.5515
Epoch 40/200, Train Loss: 28420.8021, Val Loss: 33867.7173
Epoch 60/200, Train Loss: 28273.2265, Val Loss: 33698.9863

Epoch 80/200, Train Loss: 28037.7616, Val Loss: 33429.6163
Epoch 100/200, Train Loss: 27693.5860, Val Loss: 33037.2822
Epoch 120/200, Train Loss: 27228.5033, Val Loss: 32507.8658
Epoch 140/200, Train Loss: 26636.2390, Val Loss: 31835.4775
Epoch 160/200, Train Loss: 25917.8649, Val Loss: 31018.2103
Epoch 180/200, Train Loss: 25074.7379, Val Loss: 30058.1146
Epoch 200/200, Train Loss: 24113.3227, Val Loss: 28962.0892
Deci test loss: 2803.5657
Centi test loss: 2634.3621
Milli test loss: 2940.2972
Mico test loss: 22242.8237

激活函数实验终端输出:

===== Training ReLU network =====

Epoch 20/200, Train Loss: 22255.7418, Val Loss: 26272.2201
Epoch 40/200, Train Loss: 4962.2956, Val Loss: 5028.1852
Epoch 60/200, Train Loss: 3921.5869, Val Loss: 3327.0837
Epoch 80/200, Train Loss: 3617.8698, Val Loss: 3116.0913
Epoch 100/200, Train Loss: 3403.1146, Val Loss: 3001.1186
Epoch 120/200, Train Loss: 3249.7455, Val Loss: 2942.2565
Epoch 140/200, Train Loss: 3141.7633, Val Loss: 2934.3957
Epoch 160/200, Train Loss: 3065.2934, Val Loss: 2923.0184
Epoch 180/200, Train Loss: 3014.9695, Val Loss: 2913.5199
Epoch 200/200, Train Loss: 2982.3587, Val Loss: 2919.2571

===== Training LeakyReLU network =====

Epoch 20/200, Train Loss: 21361.9628, Val Loss: 25239.1286
Epoch 40/200, Train Loss: 4811.6525, Val Loss: 4720.0527
Epoch 60/200, Train Loss: 4006.7326, Val Loss: 3384.5378
Epoch 80/200, Train Loss: 3693.9024, Val Loss: 3173.5592
Epoch 100/200, Train Loss: 3469.6493, Val Loss: 3030.3890
Epoch 120/200, Train Loss: 3309.9038, Val Loss: 2945.4171
Epoch 140/200, Train Loss: 3184.0077, Val Loss: 2922.4818
Epoch 160/200, Train Loss: 3101.4382, Val Loss: 2927.1329
Epoch 180/200, Train Loss: 3044.7196, Val Loss: 2923.4595
Epoch 200/200, Train Loss: 3004.3241, Val Loss: 2906.7160

===== Training Sigmoid network =====

Epoch 20/200, Train Loss: 26568.3883, Val Loss: 31712.5793
Epoch 40/200, Train Loss: 24945.3718, Val Loss: 29868.6338
Epoch 60/200, Train Loss: 23480.9096, Val Loss: 28186.7689
Epoch 80/200, Train Loss: 21993.7232, Val Loss: 26477.4040
Epoch 100/200, Train Loss: 20619.1416, Val Loss: 24887.5185
Epoch 120/200, Train Loss: 19242.5289, Val Loss: 23280.2169
Epoch 140/200, Train Loss: 18033.9011, Val Loss: 21860.6975
Epoch 160/200, Train Loss: 16924.8538, Val Loss: 20552.1122
Epoch 180/200, Train Loss: 15900.4869, Val Loss: 19333.5724
Epoch 200/200, Train Loss: 14953.1008, Val Loss: 18192.9443

===== Training Tanh network =====

Epoch 20/200, Train Loss: 25773.6389, Val Loss: 30798.9861
Epoch 40/200, Train Loss: 23950.4652, Val Loss: 28726.0683
Epoch 60/200, Train Loss: 22375.3980, Val Loss: 26917.1546
Epoch 80/200, Train Loss: 20938.3812, Val Loss: 25255.9906
Epoch 100/200, Train Loss: 19610.9654, Val Loss: 23713.1765
Epoch 120/200, Train Loss: 18380.8528, Val Loss: 22271.5819
Epoch 140/200, Train Loss: 17237.5195, Val Loss: 20920.0040
Epoch 160/200, Train Loss: 16172.4628, Val Loss: 19655.6430
Epoch 180/200, Train Loss: 15183.5051, Val Loss: 18470.6112
Epoch 200/200, Train Loss: 14267.9869, Val Loss: 17360.9604

===== Training Swish network =====

Epoch 20/200, Train Loss: 20115.6476, Val Loss: 23691.1254
Epoch 40/200, Train Loss: 4497.4005, Val Loss: 4184.7261
Epoch 60/200, Train Loss: 3875.4720, Val Loss: 3282.7839
Epoch 80/200, Train Loss: 3572.2693, Val Loss: 3085.5982
Epoch 100/200, Train Loss: 3352.6357, Val Loss: 2961.4223
Epoch 120/200, Train Loss: 3200.7158, Val Loss: 2919.3835
Epoch 140/200, Train Loss: 3099.0955, Val Loss: 2923.7203
Epoch 160/200, Train Loss: 3025.4077, Val Loss: 2907.5080
Epoch 180/200, Train Loss: 2983.2141, Val Loss: 2923.6613
Epoch 200/200, Train Loss: 2954.4425, Val Loss: 2931.0662
ReLU test loss: 2910.4995
LeakyReLU test loss: 2935.8343

Sigmoid test loss: 13317.2608

Tanh test loss: 12662.0868

Swish test loss: 2917.5101